

Introduction to Deep learning: an historical perspective

Stéphane Canu, INSA Rouen – Normandy University

github.com/StephaneCanu/Deep_learning_lecture



Shanghai Jiao Tong University

October 31

github.com/StephaneCanu/Deep_learning_lecture

https://github.com/StephaneCanu/Deep_learning_lecture

EDT_Mastere PythonMatlab Traducteur Adèle Anna Anna_fact Mastere Journal_Rank dico_GB_Fr Books CNU Scholar Mood

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

Deep learning lecture Edit

Manage topics

50 commits 1 branch 0 releases 1 contributor

Branch: master New pull request Create new file Upload files Find file Clone or download

StephaneCanu	Update README.md	Latest commit 7172e37 a minute ago
jupyter_notebooks	Rename TP1_MNIST_Optim.ipynb to jupyter_notebooks/TP1_MNIST_Optim.ipynb	21 hours ago
00_Main_Deep_2018.pdf	Add files via upload	2 months ago
AFIA_AG_2018.pdf	Add files via upload	20 days ago
Intro_Deep_SJTU_2018.pdf	Add files via upload	5 minutes ago
README.md	Update README.md	a minute ago
TP_Deep_1_MNIST.py	Add files via upload	2 months ago
TP_Deep_2_webcam.py	Add files via upload	2 months ago
TP_Deep_3_fine_tuning.py	Add files via upload	2 months ago
test_cheese.zip	Add files via upload	2 months ago
train_cheese.zip	Add files via upload	2 months ago

README.md

Deep learning: an historical perspective

Description

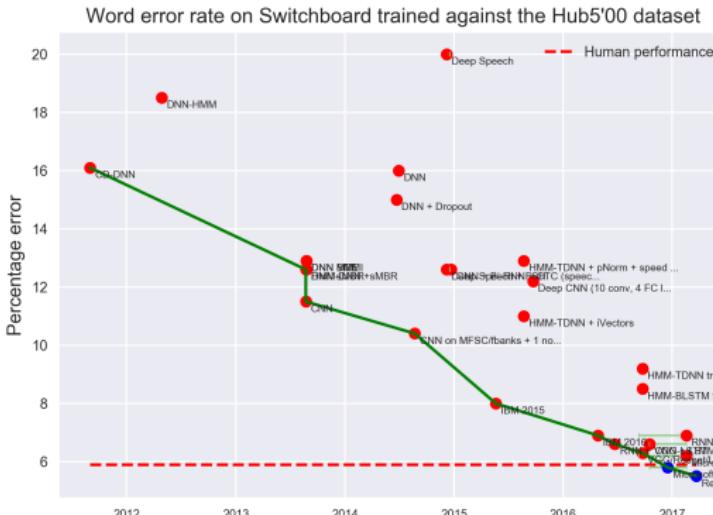
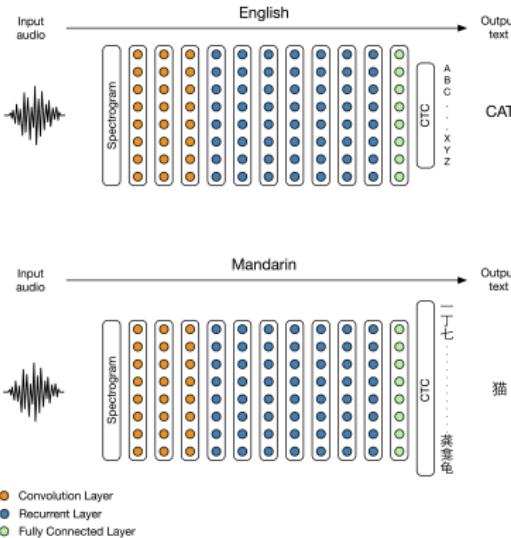
This repository contains the slides of the seminar "Deep learning: an historical perspective", given October 31 2pm-3pm at SJTU co-organized by SJTU School of Mathematical Sciences and SJTU-ParisTech Elite Institute of Technology ([Intro_Deep_SJTU_2018.pdf](#), 20.4 MB)

Road map

- ① Why deep learning?
- ② The first stage: 1890 - 1969
- ③ The second stage: 1985 - 1995
- ④ The third stage: 2006 - (2012) - 2018...
- ⑤ What's new in deep learning?
 - Big is beautiful
 - Two Hot topics: data and architecture
- ⑥ Conclusion



Deep learning for turning text into speech (and vice versa)



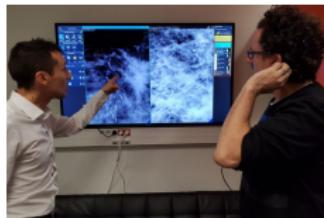
Baidu deep speech 2 (2015) and Deep voice (2017)

Trained on 9,400 hours of labeled audio with 11 million utterances.

Deep learning for healthcare



Skin cancer classification
130 000 training images
validation error rate : 28 % (human 34 %)



the Digital Mammography DREAM Challenge
640 000 mammographies (1209 participants)
5 % less false positive

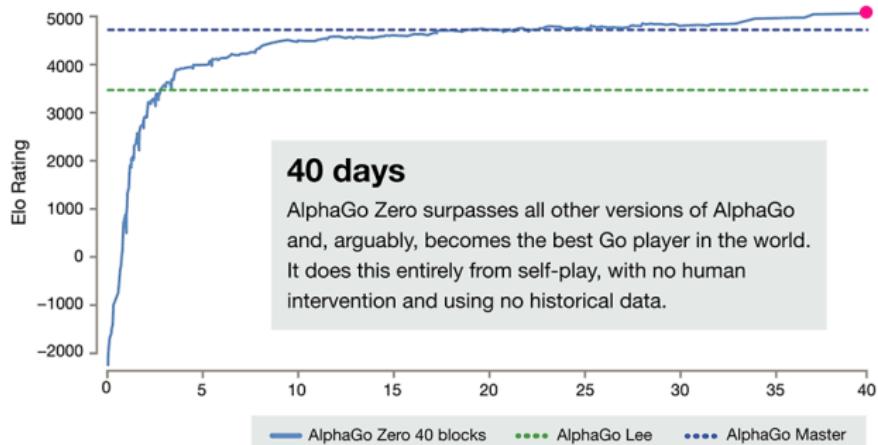


heart rate analysis
500 000 ECG
precision 92.6 % (humain 80.0 %) sensitivity 97 %

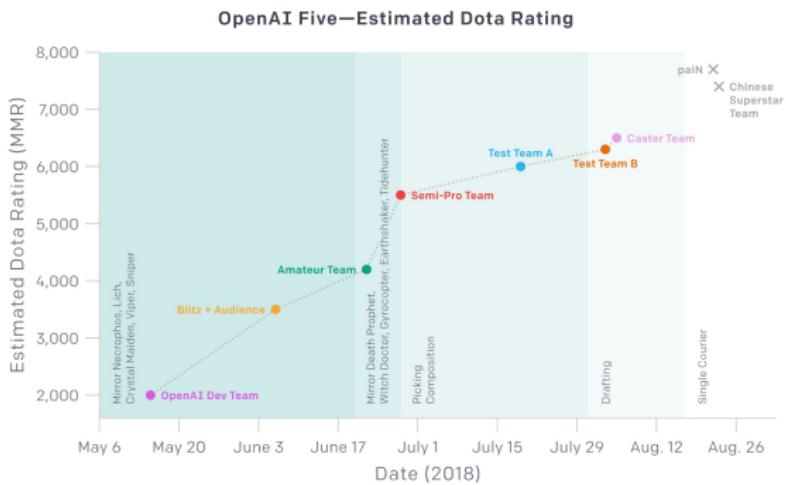
Statistical machine learning: retrieving correlations

with deep learning end-to-end architecture
"April showers bring May flowers"

Deep learning success in playing GO

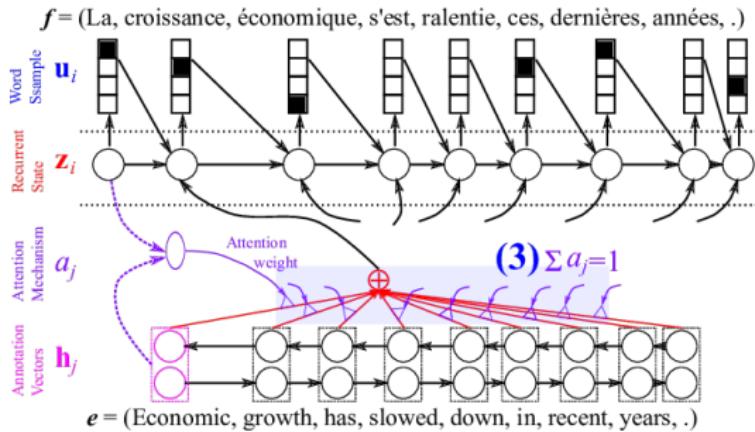


Deep learning progresses in playing Dota 2



- separate LSTM for each hero
- 180 years/days of games against itself
- Proximal Policy Optimization
- 256 GPUs and 128,000 CPU
- the OpenAI Five is very much still a work in progress project

Deep learning (limited) success in NLP



Learning to translate with 36 million sentences

- Near Human-Level Performance in Grammatical Error Correction
- Achieving Human Parity on Automatic News Translation

Deep learning to drive: the Rouen autonomous lab



Driving Video Database = 100.000 videos – 120 million images

- When It Comes to Safety, Autonomous Cars Are Still "Teen Drivers"
- companies are developing many different levels of automation

So far, so good

- Deep learning performance breakthrough
 - ▶ Low level perception tasks: speech, image and video processing, natural language processing, games...
 - ▶ ...and specific tasks in health care, astronomy...
- It requires
 - ▶ Big data
 - ▶ Big computers
 - ▶ Specific tasks

The screenshot shows a news article from Fortune magazine. At the top right is a 'SUBSCRIBE' button. Below it is the word 'FORTUNE' next to a three-line menu icon. The author's name, 'By Roger Parloff', is followed by a small note: 'Illustration by Justin Metz'. The date 'SEPTEMBER 28, 2016, 5:00 PM EDT' is also present. The main title of the article is 'WHY DEEP LEARNING IS SUDDENLY CHANGING YOUR LIFE', displayed prominently in white text on a red background.

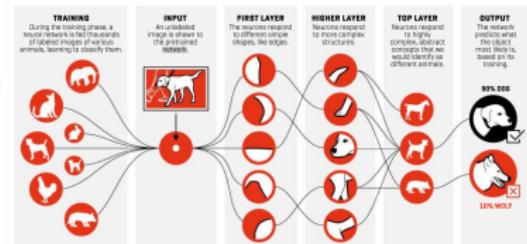
Decades-old discoveries are now electrifying the computing industry and will soon transform corporate America.

Over the past four years, readers have doubtlessly noticed quantum leaps in the quality of a wide range of everyday technologies.

Most obviously, the speech-recognition functions on our smartphones work much better than they used to. When we use a voice command to call our spouses, we reach them now. We aren't connected to Amtrak or an angry ex.

Road map

- 1 Why deep learning?
- 2 The first stage: 1890 - 1969
- 3 The second stage: 1985 - 1995
- 4 The third stage: 2006 - (2012) - 2018...
- 5 What's new in deep learning?
 - Big is beautiful
 - Two Hot topics: data and architecture
- 6 Conclusion



The neural networks time line

- The first stage: 1890 - 1969

~1890 Ramón y Cajal: the biological neuron

1943 McCulloch & Pitts formal neuron

1949 Hebb's rule

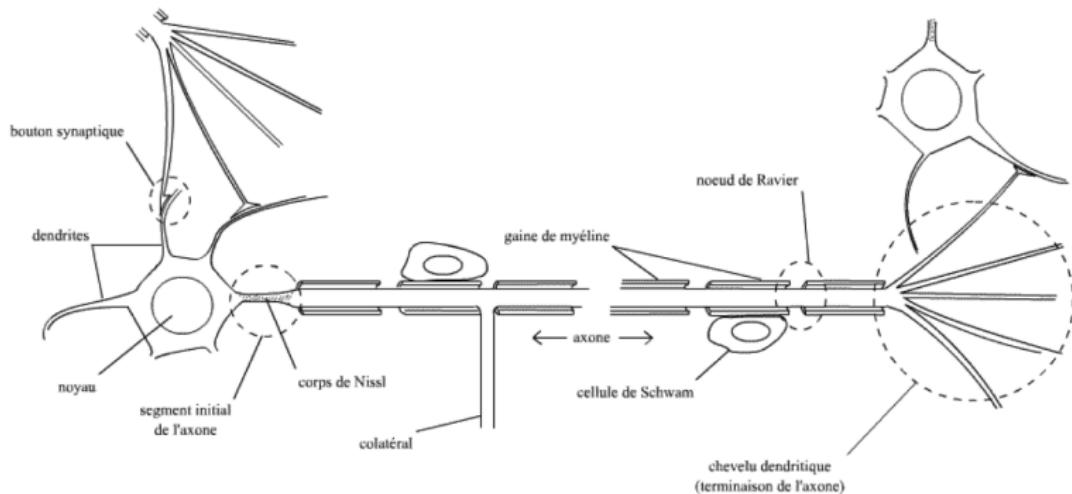
1958 Rosenblatt's Perceptron: learning with stochastic gradient

1969 Minsky & Papert: stop – the 1st NN winter

- The second stage: 1985 - 1995

- The third stage: 2006 - (2012) - 2018...

The biological neuron



The neural networks time line

- The first stage: 1890 - 1969

~1890 Ramón y Cajal: the biological neuron

1943 McCulloch & Pitts formal neuron

1949 Hebb's rule

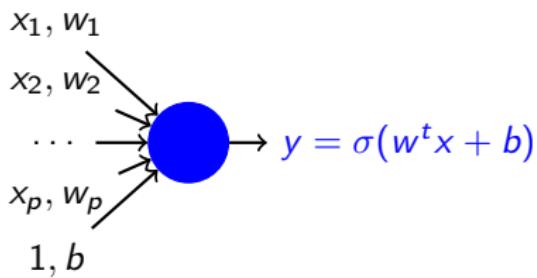
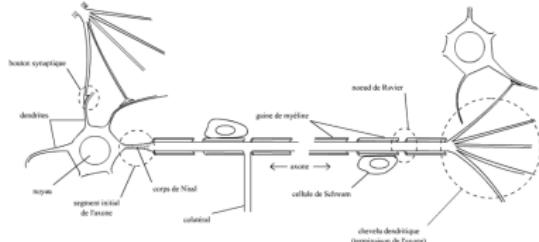
1958 Rosenblatt's Perceptron: learning with stochastic gradient

1969 Minsky & Papert: stop – the 1st NN winter

- The second stage: 1985 - 1995

- The third stage: 2006 - (2012) - 2018...

McCulloch & Pitts formal neuron 1943



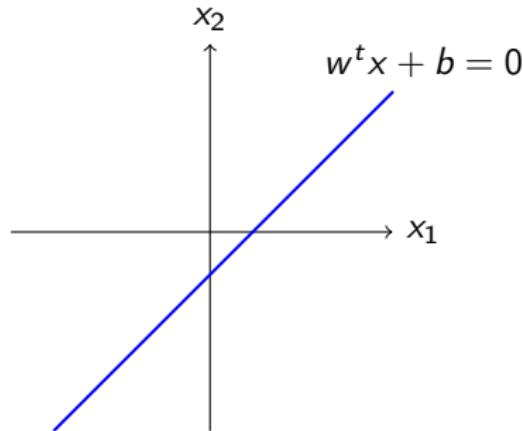
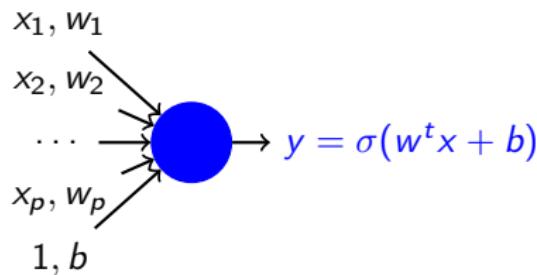
x input $\in \mathbb{R}^p$

w weight, b bias

σ activation function

y output $\in \mathbb{R}$

The artificial neuron as a linear threshold unit



x input $\in \mathbb{R}^p$

σ activation function (non linear)

w weight, b bias

$\mathbb{R} \mapsto \mathbb{R}$

a activation, $a = w^t x + b$

$a \rightarrow y = \sigma(a)$

σ activation function

Φ transfer function

Φ transfer function

y output $\in \mathbb{R}$

$\mathbb{R}^p \mapsto \mathbb{R}$

$x \rightarrow y = \Phi(x) = \sigma(w^t x + b)$

The neural networks time line

- The first stage: 1890 - 1969

~1890 Ramón y Cajal: the biological neuron

1943 McCulloch & Pitts formal neuron

1949 Hebb's rule

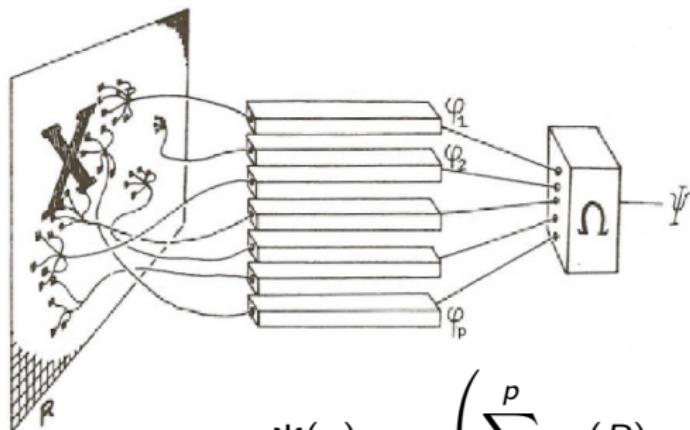
1958 Rosenblatt's Perceptron: learning with stochastic gradient

1969 Minsky & Papert: stop – the 1st NN winter

- The second stage: 1985 - 1995

- The third stage: 2006 - (2012) - 2018...

The formal neuron as a learning machine: fit the w



$$\Psi(x) = \sigma \left(\sum_{j=1}^p \varphi_j(R) w_j + b \right)$$

Rosenblatt's Perceptron, 1958 (Widrow & Hoff's Adaline, 1960)

given n pairs of input-output data $x_i = \varphi_j(R_i), t_i, i = 1, n$

find w such that

$$\underbrace{\sigma(w^t x_i)}_{\text{prediction of the model}} = \underbrace{t_i}_{\text{ground truth}}$$

Cost minimization (energy-based model)

Minimize a loss $\min_{w \in \mathbb{R}^{p+1}} \sum_{i=1}^n \text{loss}(w) \quad \text{loss}(w) = (\sigma(w^t x_i) - t_i)^2$

Gradient descent $w \leftarrow w - \rho d$ $d = \sum_{i=1}^n \nabla_w \text{loss}(w)$

Stochastic gradient $d = \nabla_w \text{loss}(w)$

Algorithm 1 Gradient epoch

Data: w initialization, ρ stepsize

Result: w

for $i=1, n$ **do**

$x_i, t_i \leftarrow$ pick a point i

$d \leftarrow d + \nabla_w \text{loss}(w, x_i, t_i)$

end

$w \leftarrow w - \rho d$

Algorithm 2 Stochastic gradient

Data: w initialization, ρ stepsize

Result: w

for $i=1, n$ **do**

$x_i, t_i \leftarrow$ pick a point i

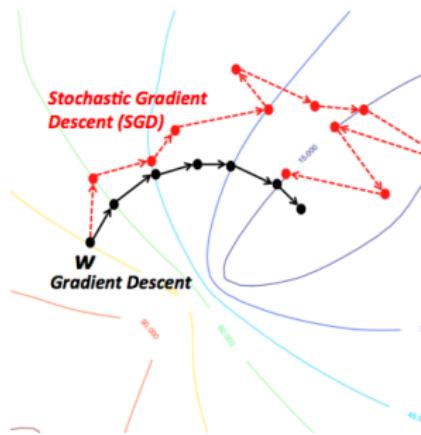
$d \leftarrow \nabla_w \text{loss}(w, x_i, t_i)$

$w \leftarrow w - \rho d$

end

Accelerating the stochastic gradient

- stochastic average (mini batch)
 - ▶ parameters (Polyak and Juditsky, 1992)
 - ▶ gradients SAG-A, (Le Roux et al 2012)
 - ▶ variance reduction (Johnson, Zhang, 2013)
- convergence acceleration
 - ▶ Nesterov's method (1983)
 - ▶ momentum (heuristic)
- acceleration and averaging
 - ▶ (Dieuleveut, Flammarion & Bach, 2016)
- stepsize adaptation
 - ▶ RMSprop (Tieleman & Hinton, 2012)
 - ▶ Adaptive Moment Estimation – ADAM (Kingma & Ba, 2015)
 - ▶ AMSGRAD (Reddi et al, BPA ICRL 2018)

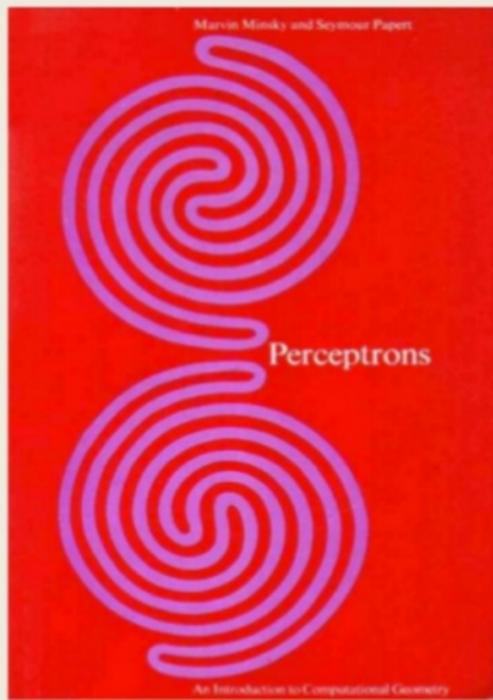


The neural networks time line

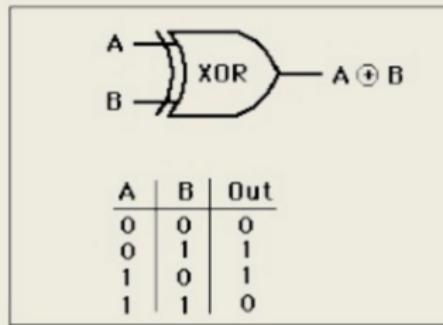
- The first stage: 1890 - 1969
 - ~1890 Ramón y Cajal: the biological neuron
 - 1943 McCulloch & Pitts formal neuron
 - 1949 Hebb's rule
 - 1958 Rosenblatt's Perceptron: learning with stochastic gradient
 - 1969 Minsky & Papert: stop – the 1st NN winter
- The second stage: 1985 - 1995
- The third stage: 2006 - (2012) - 2018...

However, linear neurons ar linear

1969: Perceptrons can't do XOR!



<http://www.i-programmer.info/images/stories/BabBag/AI/book.jpg>



<http://hyperphysics.phy-astr.gsu.edu/hbase/electronic/ietron/xor.gif>



Minsky & Papert

<https://constructingkids.files.wordpress.com/2013/05/minsky-papert-71-csolomon-x640.jpg>

The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995

1985 Rumelhart, Hinton & Williams; Le Cun: go - backpropagation

1989 Universal Approximation Cybenko-Hornik-Funahashi Theorem

1989 Y. Le Cun's convolutional neural networks

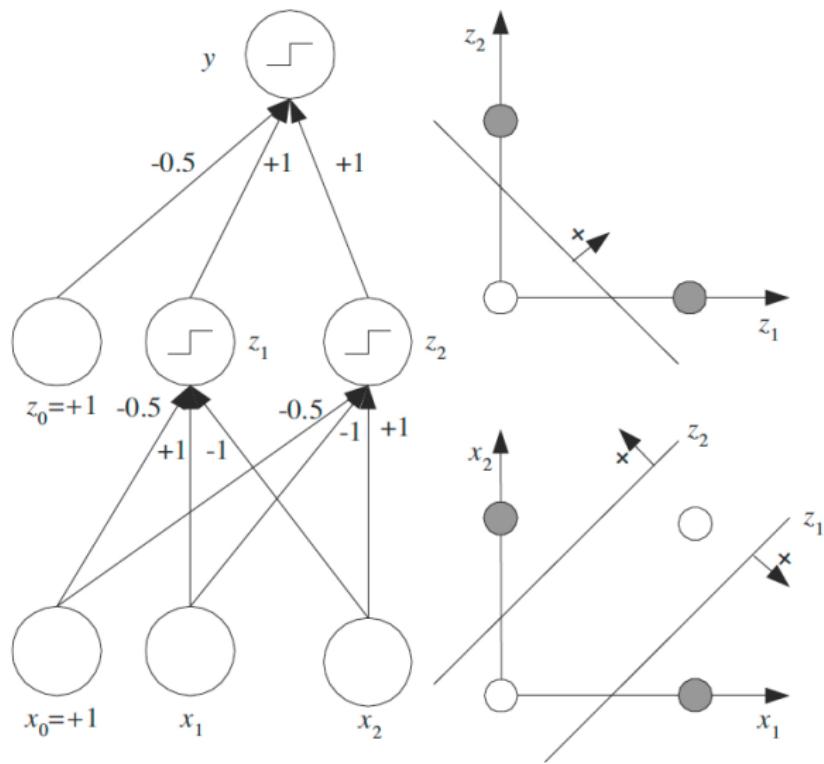
1995 Recurrent neural networks, LSTM

1995 SVM

2004 Caltech 101: the 2nd NN winter

- The third stage: 2006 - (2012) - 2018...

Non linearity combining linear neurons: the Xor case



Neural networks

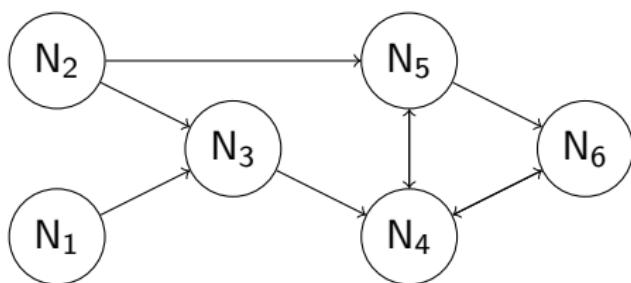
Definition: Neural network

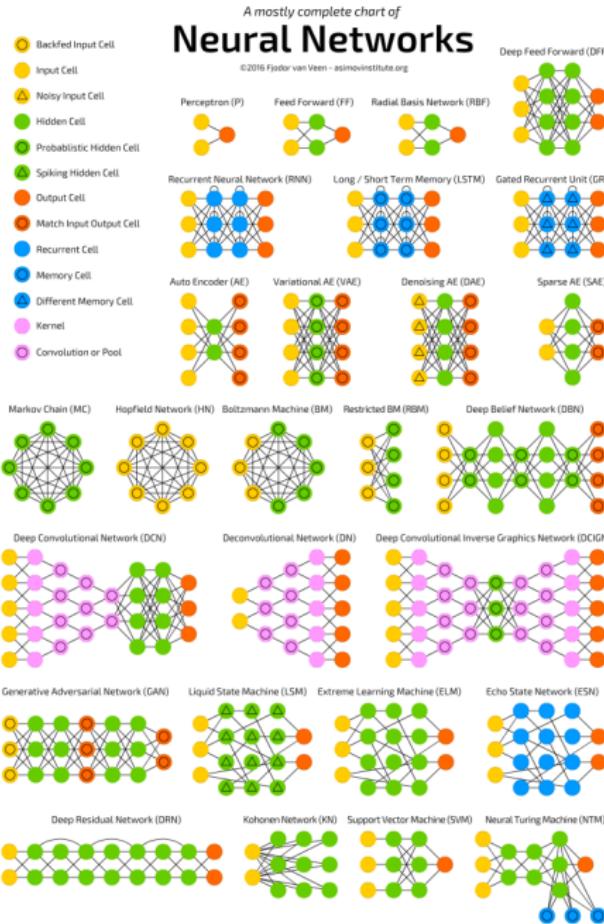
A neural network is an oriented graph of formal neurons

When two neurons are connected (linked by an oriented edge of the graph), the output of the head neuron is used as an input by the tail neuron. It can be seen as a weighted directed graph

3 different neurons are considered:

- input neurons (connected with the input)
- output neurons
- hidden neurons





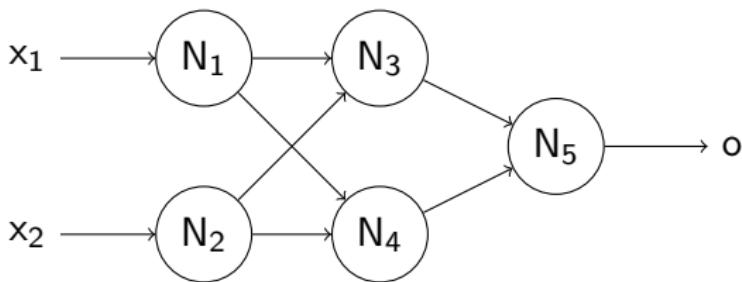
The Multiplayred peceptron (MLP)

Definition: Multiplayred peceptron

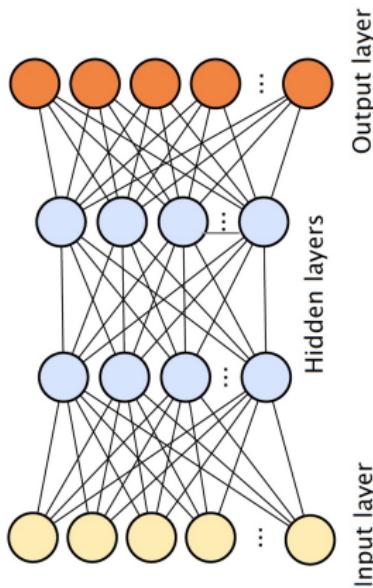
A Multiplayred peceptron is an acyclic neural network,

where the neurons are structued in successive layers, begining by an input layer and finishing with an output layer.

Example: The X-or neural network is a MLP with a single hidden unit with 2 hidden neurons.



MLP training with back propagation (and SGD)



$$y = \sigma(W_3 h^{(2)})$$

$$\nabla_{W_3} J = (y - y_a) \sigma'(W_3 h^{(2)}) h^{(2)}$$

↑

↓

$$h^{(2)} = \sigma(W_2 h^{(1)})$$

$$\nabla_{W_2} J = \nabla_{h^{(2)}} J \sigma'(W_2 h^{(1)}) h^{(1)}$$

↑

↓

$$h^{(1)} = \sigma(W_1 x)$$

$$\nabla_{W_1} J = \nabla_{h^{(1)}} J \sigma'(W_1 x) x^\top$$

↑

x

$$y = \sigma(W_3 \sigma(W_2 \sigma(W_1 x)))$$

backpropagation = chain rule (autodiff)

Used to learn internal representation W_1, W_2, W_3

Back propagation is differential learning



Yann LeCun

5 janvier ·

OK, Deep Learning has outlived its usefulness as a buzz-phrase.
Deep Learning est mort. Vive Differentiable Programming!

Numpy

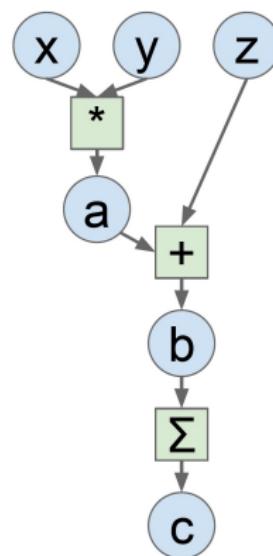
```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995
 - 1985 Rumelhart, Hinton & Williams; Le Cun: go - backpropagation
 - 1989 Universal Approximation Cybenko-Hornik-Funahashi Theorem
 - 1989 Y. Le Cun's convolutional neural networks
 - 1995 Recurrent neural networks, LSTM
 - 1995 SVM
 - 2004 Caltech 101: the 2nd NN winter
- The third stage: 2006 - (2012) - 2018...

MLP with one hidden layer as universal approximator

Universal approximation theorem for MLP

- given any $\varepsilon > 0$
- for any continuous function f on compact subsets of \mathbb{R}^p
- for any admissible activation function σ (not a polynomial)
- there exists h , $W_1 \in \mathbb{R}^{p \times h}$, $b \in \mathbb{R}^h$, $c \in \mathbb{R}$ and $w_2 \in \mathbb{R}^h$ such that

$$\|f(x) - w_2\sigma(W_1x + b) + c\|_\infty \leq \varepsilon$$

SVM, Boosting and Random Forest also

The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995

1985 Rumelhart, Hinton & Williams; Le Cun: go - backpropagation

1989 Universal Approximation Cybenko-Hornik-Funahashi Theorem

1989 Y. Le Cun's convolutional neural networks

1995 Recurrent neural networks, LSTM

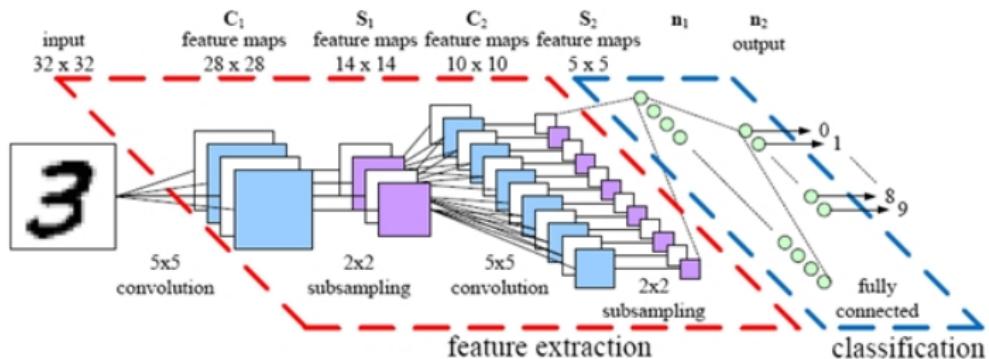
1995 SVM

2004 Caltech 101: the 2nd NN winter

- The third stage: 2006 - (2012) - 2018...

OCR: MNIST database (LeCun, 1989)

3	8	6	9	6	4	5	3	8	4	5	2	3	8	4	8
1	5	0	5	9	7	4	1	0	3	0	6	2	9	9	4
1	3	6	8	0	7	7	6	8	9	0	3	8	3	7	7
8	4	4	1	2	9	8	1	1	0	6	6	5	0	1	1
7	2	7	3	1	4	0	5	0	6	8	7	6	8	9	9
4	0	6	1	9	2	6	3	1	4	4	5	6	6	1	7
2	8	6	9	7	0	9	1	6	2	8	3	6	4	9	5
8	6	8	7	8	8	6	9	1	7	6	0	9	6	7	0



use convolution layers

The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995

1985 Rumelhart, Hinton & Williams; Le Cun: go - backpropagation

1989 Universal Approximation Cybenko-Hornik-Funahashi Theorem

1989 Y. Le Cun's convolutional neural networks

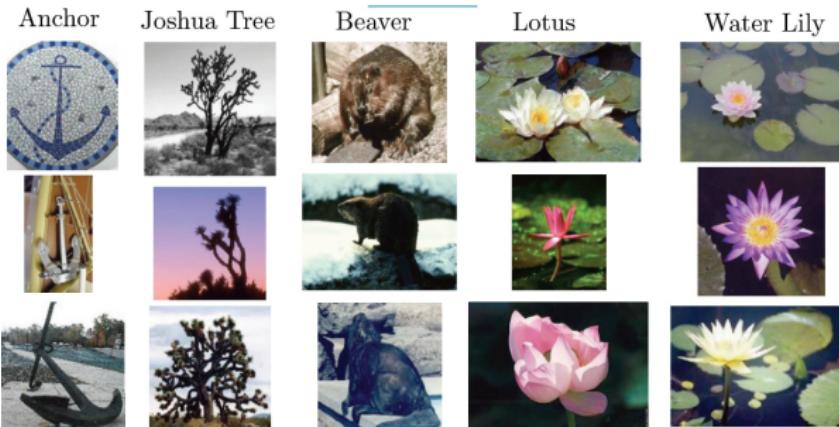
1995 Recurrent neural networks, LSTM

1995 SVM

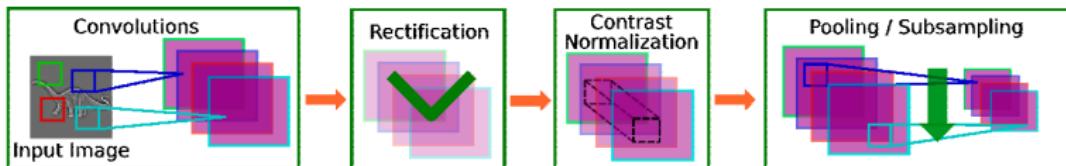
2004 Caltech 101: the NN winter

- The third stage: 2006 - (2012) - 2018...

The caltech 101 database (2004)



- 101 classes,
- 30 training images per category
- ...and the winner is NOT a deep network
 - ▶ dataset is too small



use convolution + Rectification + Normalization + Pooling

The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995
- The third stage: 2006 - (2012) - 2018...

- 2004 CIFAR Research Program: Learning in Machines & Brains
- 2006 Deep learning: Bengio's, Hinton's RBM, Y LeCun's proposals
- 2010 Andrew Ng's GPU for Deep GPU
- 2011 Deep frameworks, tools (theano, torch, cuda-convnet...)
- 2012 **ImageNet – AlexNet**
- 2013 M. Zuckerberg at NIPS the deep fashion
- 2014 Representation learning fine tuning
- 2015 Deep learning in the industry: speech, translation, image...
- 2016 Goodfellow's generative adversarial networks (GAN)
- 2017 Reinforcement learning: DeepMind's GO
- 2018 Automatic design, adversarial defense, green learning, theory...

The image net database (Deng et al., 2012)



ImageNet = 15 million high-resolution images of 22,000 categories.
Large-Scale Visual Recognition Challenge (a subset of ImageNet)

- 1000 categories.
- 1.2 million training images,
- 50,000 validation images,
- 150,000 testing images.

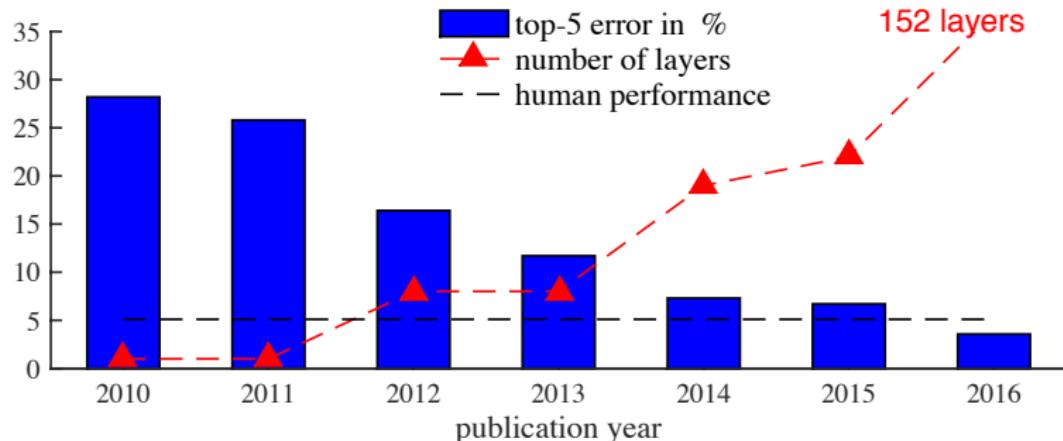
A new fashion in image processing

2012 Teams	%error	2013 Teams	%error	2014 Teams	%error
Supervision (Toronto)	15.3	Clarifai (NYU spinoff)	11.7	GoogLeNet	6.6
ISI (Tokyo)	26.1	NUS (singapore)	12.9	VGG (Oxford)	7.3
VGG (Oxford)	26.9	Zeiler-Fergus (NYU)	13.5	MSRA	8.0
XRCE/INRIA	27.0	A. Howard	13.5	A. Howard	8.1
UvA (Amsterdam)	29.6	OverFeat (NYU)	14.1	DeeperVision	9.5
INRIA/LEAR	33.4	UvA (Amsterdam)	14.2	NUS-BST	9.7
		Adobe	15.2	TTIC-ECP	10.2
		VGG (Oxford)	15.2	XYZ	11.2
		VGG (Oxford)	23.0	UvA	12.1

shallow approaches

deep learning

ImageNet results



2012 Alex Net

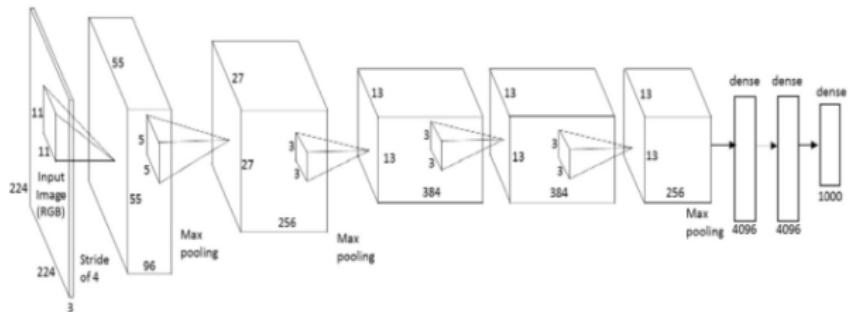
2013 ZFNet

2014 VGG

2015 GoogLeNet / Inception

2016 Residual Network

Deep architecture for image net (15%)



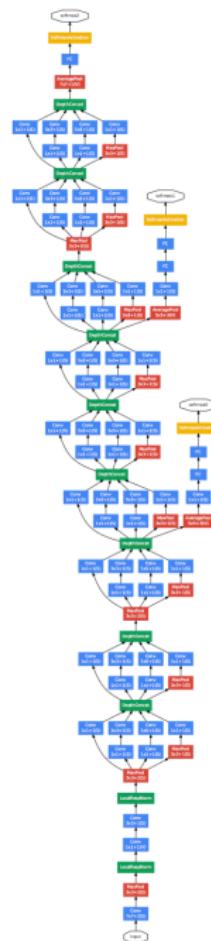
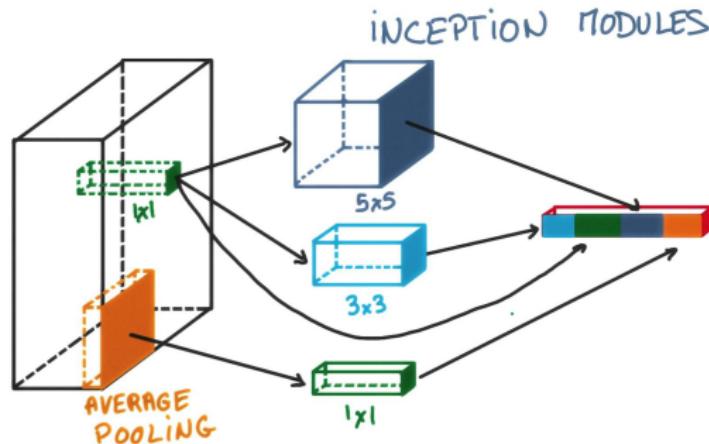
The *Alex Net* architecture [Krizhevsky, Sutskever, Hinton, 2012]

Convolution + Rectification (ReLU) + Normalization + Pooling

- 60 million parameters
- using 2 GPU – 6 days
- regularization
 - ▶ data augmentation
 - ▶ dropout
 - ▶ weight decay



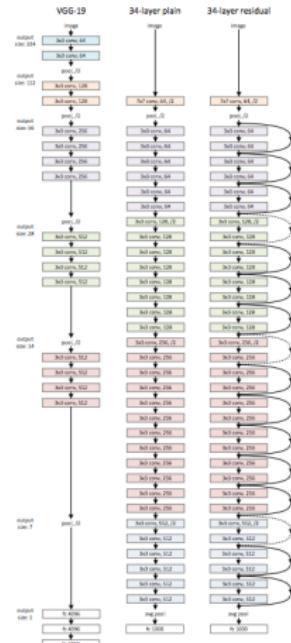
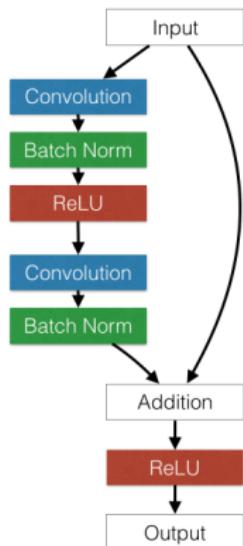
From 15% to 7%: Inceptionism



Network in a network (deep learning lecture at Udacity)

Christian Szegedy et. al. Going deeper with convolutions. CVPR 2015.

From 7% to 3%: Residual Nets



Beating the gradient vanishing effect

K. He et al., 2016

The neural networks time line

- The first stage: 1890 - 1969
- The second stage: 1985 - 1995
- The third stage: 2006 - (2012) - 2018...

- 2004 CIFAR Research Program: Learning in Machines & Brains
- 2006 Deep learning: Bengio's, Hinton's RBM, Y LeCun's proposals
- 2010 Andrew Ng's GPU for Deep GPU
- 2011 Deep frameworks, tools (theano, torch...)
- 2012 ImageNet – AlexNet
- 2013 M. Zuckerberg at NIPS: the deep fashion
- 2014 Representation learning fine tuning
- 2015 Deep learning in the industry: speech, translation, image...
- 2016 Goodfellow's generative adversarial networks (GAN)
- 2017 Reinforcement learning: DeepMind's GO
- 2018 Automatic design, adversarial defense, green learning, theory...

Deep learning, AI and the industry

*backpropagation,
boltzmann machines*



Geoff Hinton
Google

convolution



Yann Lecun
Facebook

*stacked auto-
encoders*



Yoshua Bengio
U. of Montreal

GPU utilization



Andrew Ng
Baidu

dropout



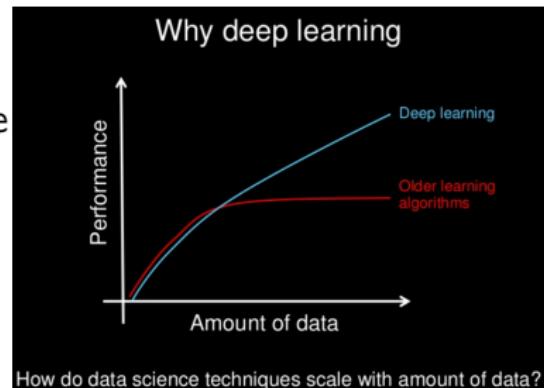
Alex Krizhevsky
Google

- data science, **artificial intelligence** and deep learning
- the BATX - GAFAM vision
 - ▶ they got the infrastructure (hard+software)
 - ▶ they got the data
 - ▶ deep learning bridges the gap between applications and ML

In 2016, Google Chief Executive Officer (CEO) Sundar Pichai said, Machine learning [a subfield of AI] is a core, transformative way by which we're rethinking how we're doing everything.

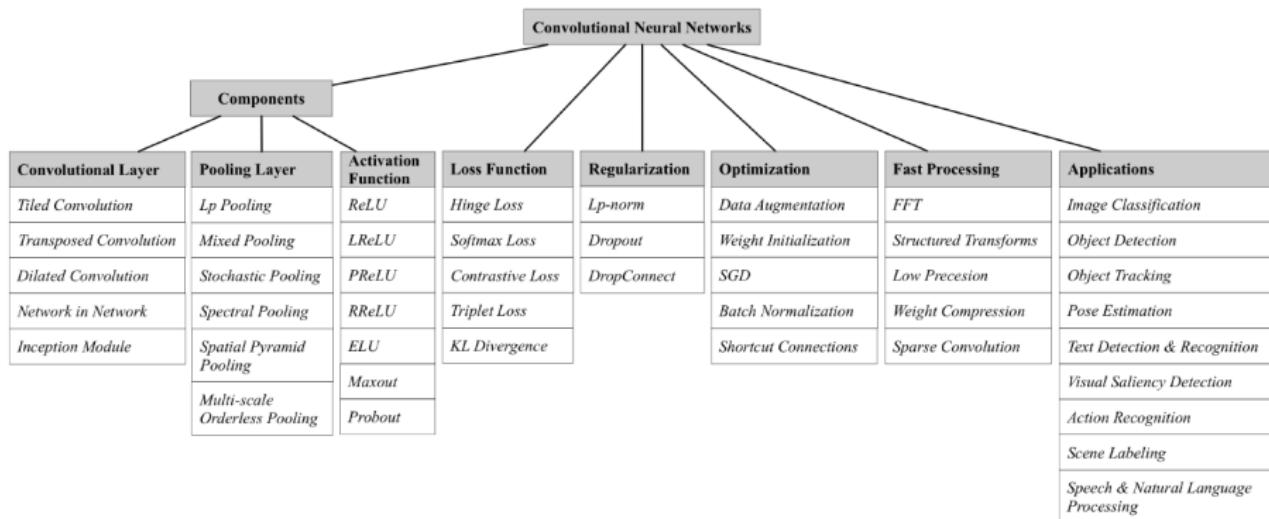
Road map

- 1 Why deep learning?
- 2 The first stage: 1890 - 1969
- 3 The second stage: 1985 - 1995
- 4 The third stage: 2006 - (2012) - 2018...
- 5 What's new in deep learning?
 - Big is beautiful
 - Two Hot topics: data and architecture
- 6 Conclusion



What's new with deep learning

- a lot of **data** (big data)
- big computing resources (**hardware & software**),
- big **model** (deep vs. shallow)
 - new architectures
 - new learning tricks

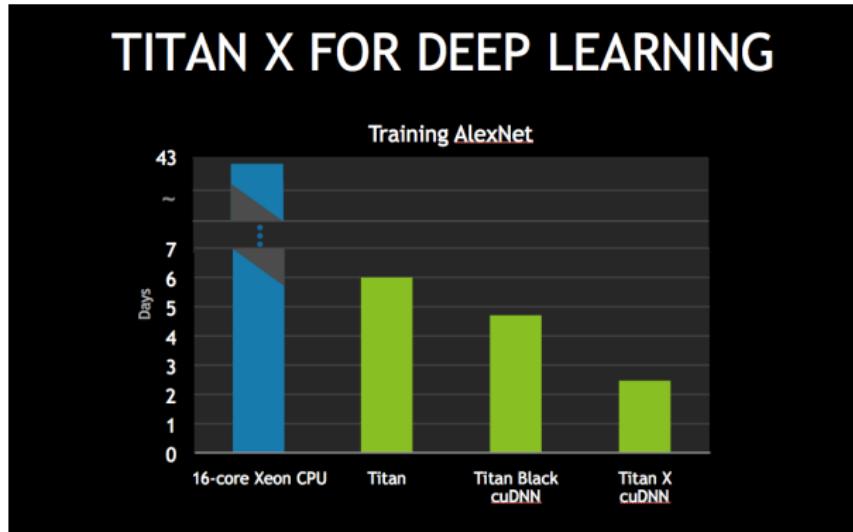


Big data: a lot of available training data



- ImageNet: 1,200,000x256x256x3 (about 200GB) block of pixels
- MS COCO for supervised learning
 - ▶ Multiple objects per image
 - ▶ More than 300,000 images
 - ▶ More than 2 Million instances
 - ▶ 80 object categories
 - ▶ 5 captions per image
- YFCC100M for unsupervised learning
- Google Open Images, 9 million URLs to images annotated over 6000 categories
- Visual genome: data + knowledge <http://visualgenome.org/>

Big computers: GPU needed



Now 2 hours with Nvidia DGX-1, and enough Memory

Yann LeCun:

learning a relevant model takes 3 weeks



Yann LeCun a partagé la publication de Yangqing Jia.
11 h · ©

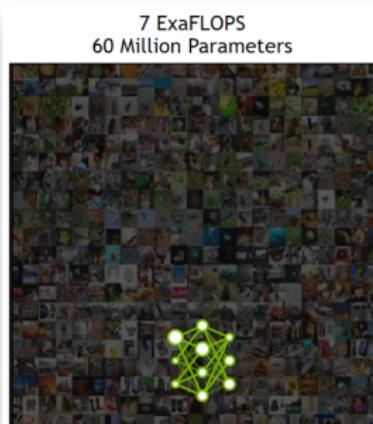
Want to train ResNet50 on ImageNet in 1 hour?
Yes (256 GPUs required).

big software: deep learning frameworks

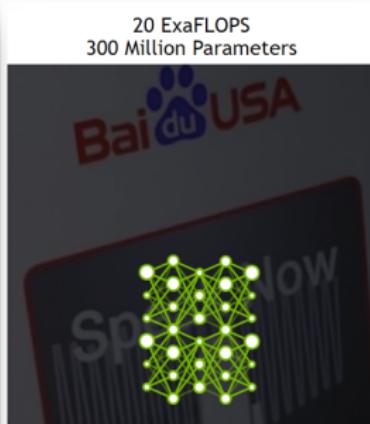
	Languages	Tutorials and training materials	CNN modeling capability	RNN modeling capability	Architecture: easy-to-use and modular front end	Speed	Multiple GPU support	Keras compatible
Theano	Python, C++	++	++	++	+	++	+	+
Tensor-Flow	Python	+++	+++	++	+++	++	++	+
Torch	Lua, Python (new)	+	+++	++	++	+++	++	
Caffe	C++	+	++		+	+	+	
MXNet	R, Python, Julia, Scala	++	++	+	++	++	+++	
Neon	Python	+	++	+	+	++	+	
CNTK	C++	+	+	+++	+	++	+	

Tensoflow (Google) is the most popular with Keras. Pytorch is a challenger.

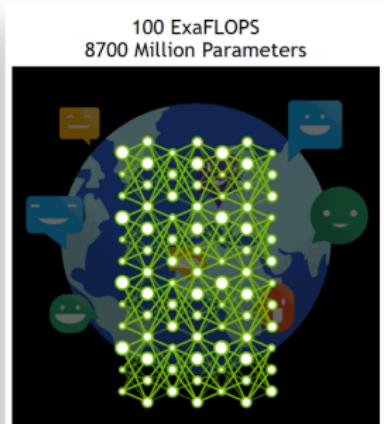
Big architectures



2015 - Microsoft ResNet
Superhuman Image Recognition



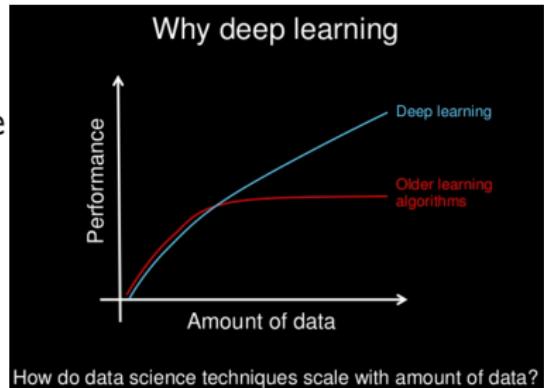
2016 - Baidu Deep Speech 2
Superhuman Voice Recognition



2017 - Google Neural Machine Translation
Near Human Language Translation

Road map

- 1 Why deep learning?
- 2 The first stage: 1890 - 1969
- 3 The second stage: 1985 - 1995
- 4 The third stage: 2006 - (2012) - 2018...
- 5 What's new in deep learning?
 - Big is beautiful
 - Two Hot topics: data and architecture
- 6 Conclusion

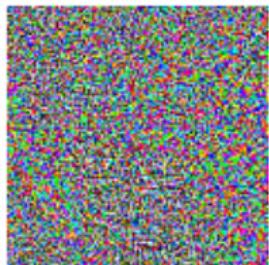


<https://www.forbes.com/sites/mariyayao/2018/02/05/12-amazing-deep-learning-breakthroughs-of-2017>

Deep neural networks are easily fooled (1/2)



$$+ .007 \times$$



=



x

“panda”

57.7% confidence

$$\text{sign}(\nabla_x J(\theta, x, y))$$

“nematode”

8.2% confidence

$$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$$

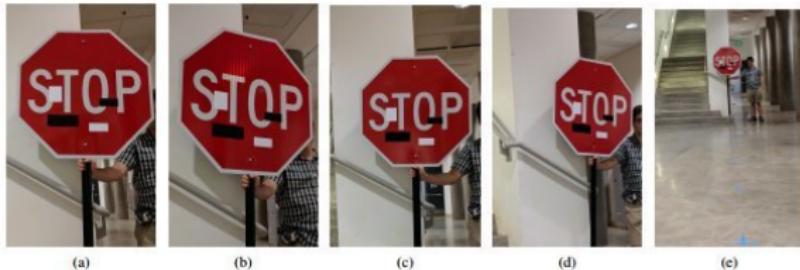
“gibbon”

99.3 % confidence

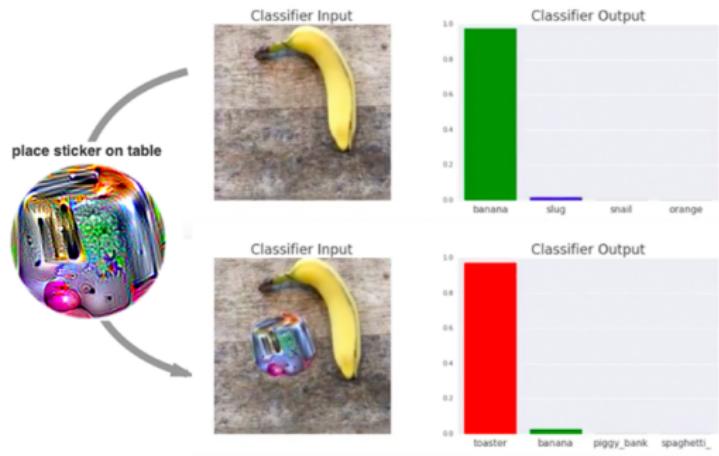
Explaining and Harnessing Adversarial Examples, Ian J. Goodfellow, Jonathon Shlens, Christian Szegedy, 2015

<https://arxiv.org/abs/1412.6572>

Adversarial examples (2/2)



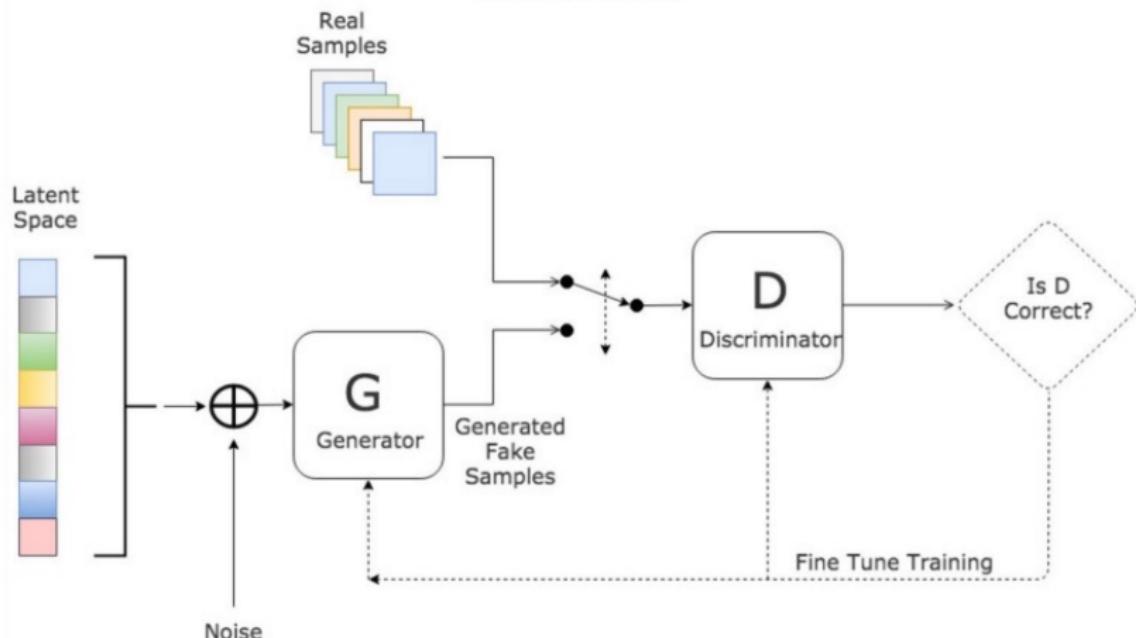
Adversarial Examples for Evaluating Reading Comprehension Systems, Robin Jia, Percy Liang, 2017



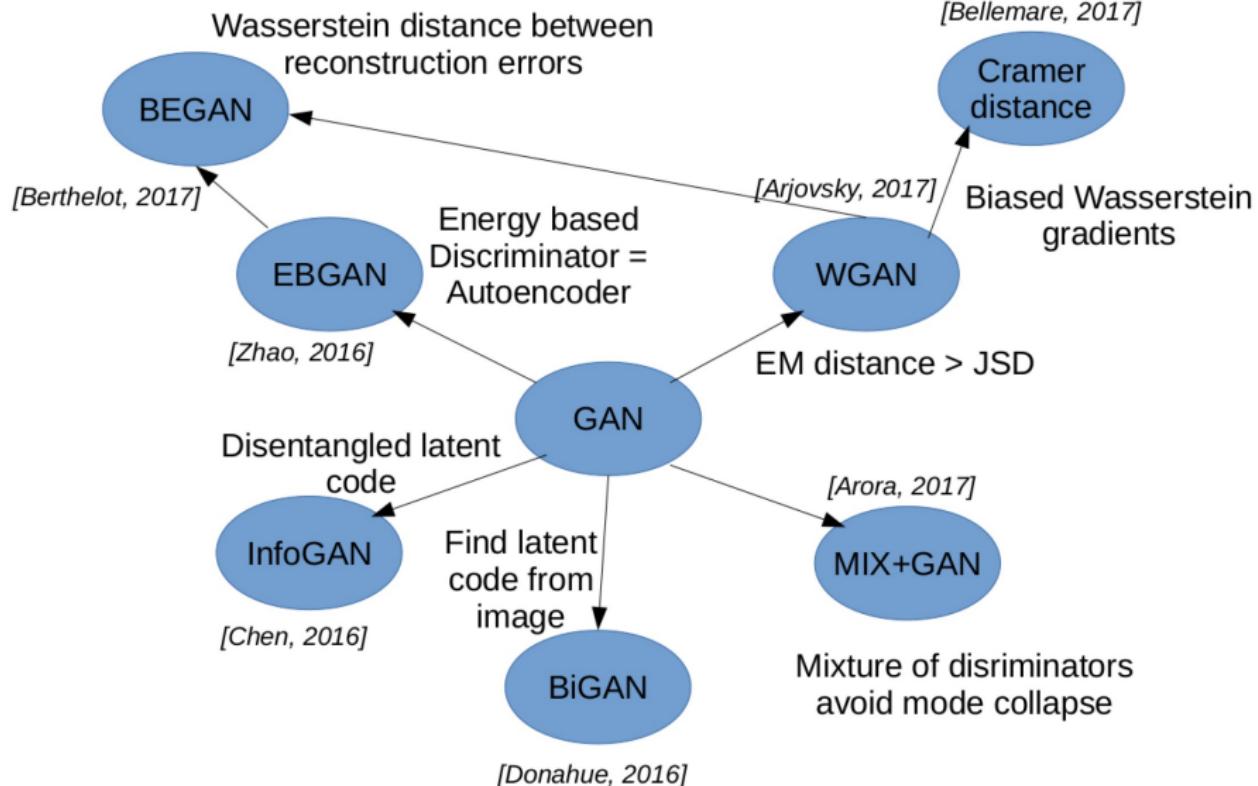
Adversarial Patch Tom B. Brown, Dandelion Mané, Aurko Roy, Martin Abadi, Justin Gilmer, 2017

Generative models

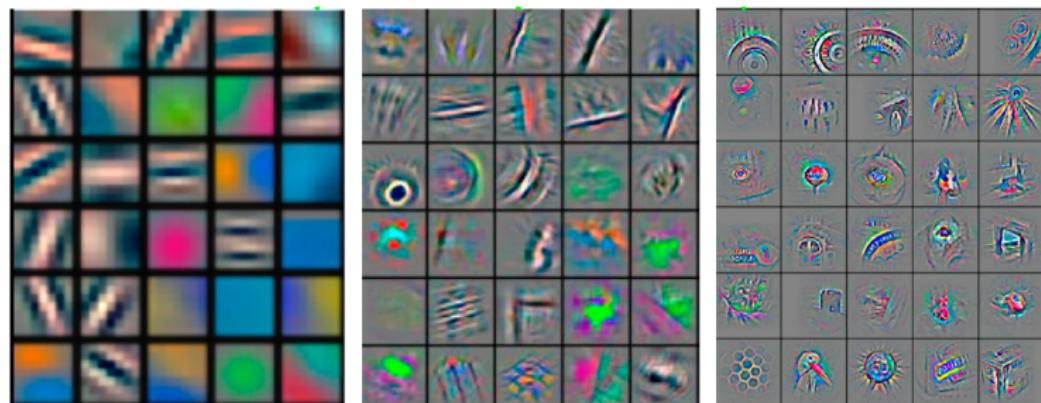
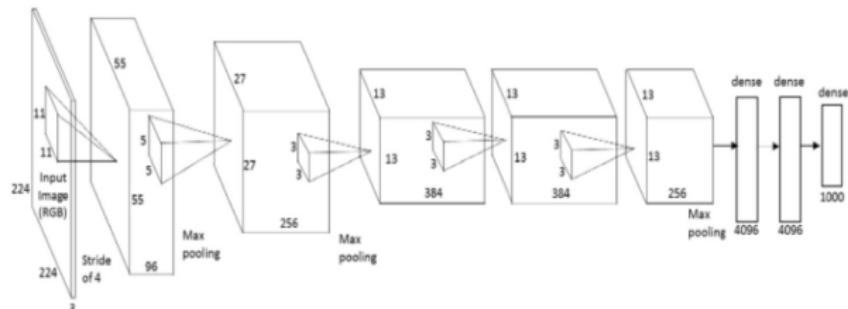
Generative Adversarial Network



Other Generative architectures

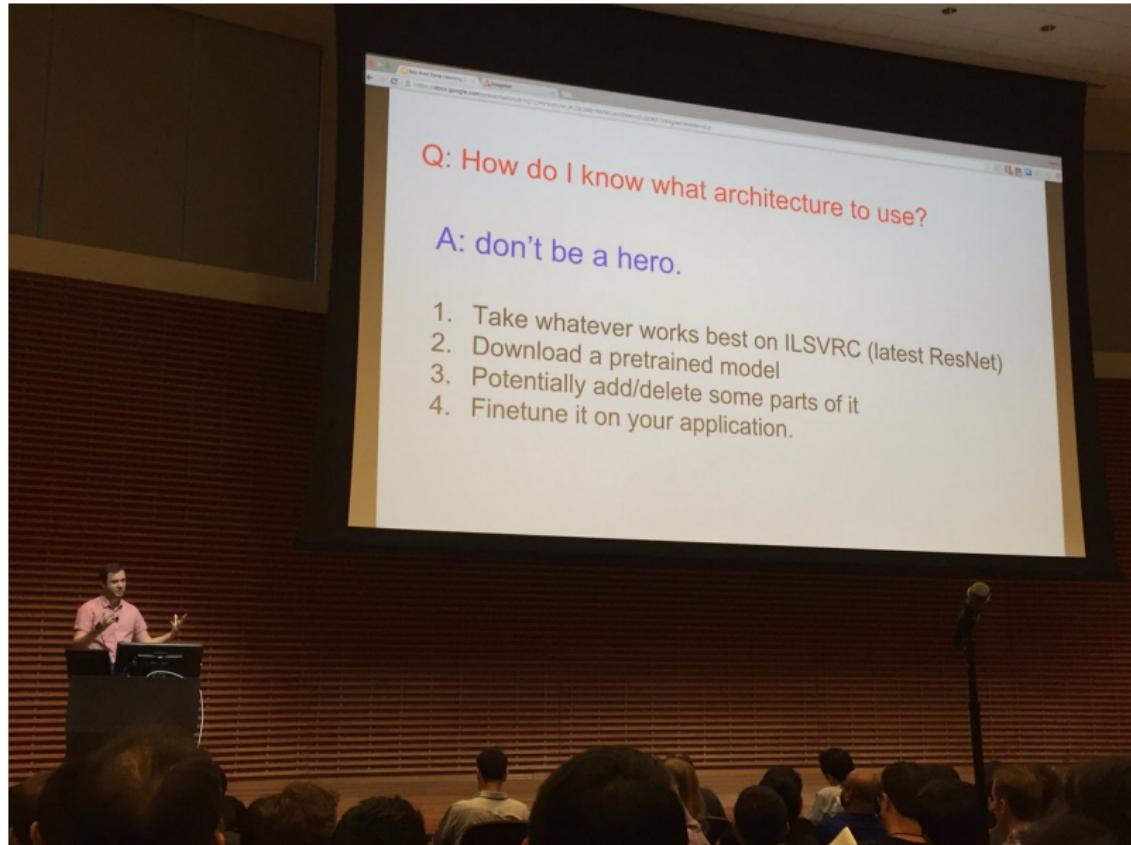


AlexNet works because of learning internal representation



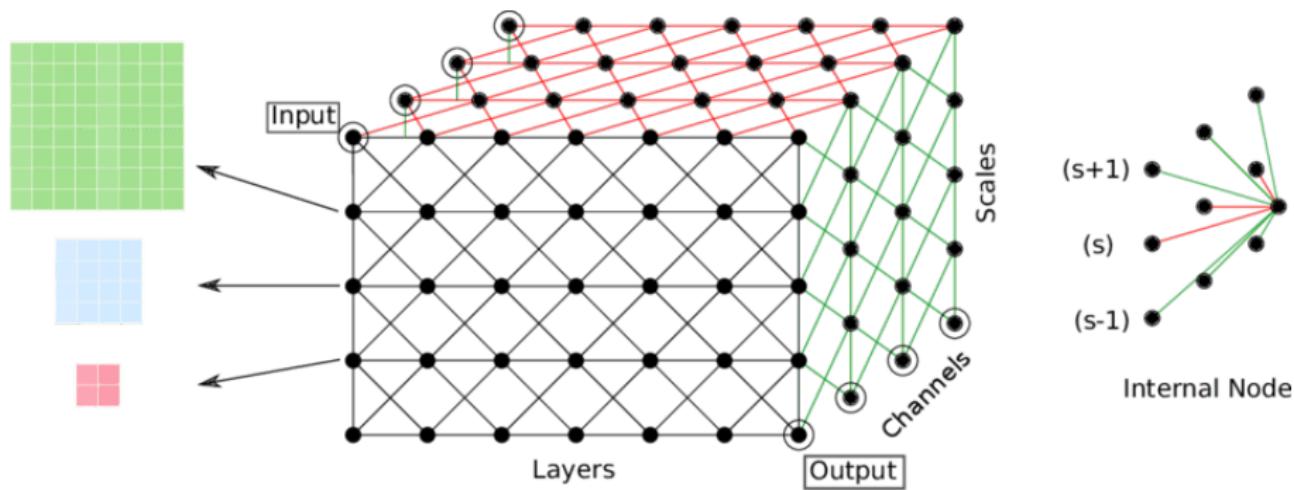
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

How to start with deep learning?

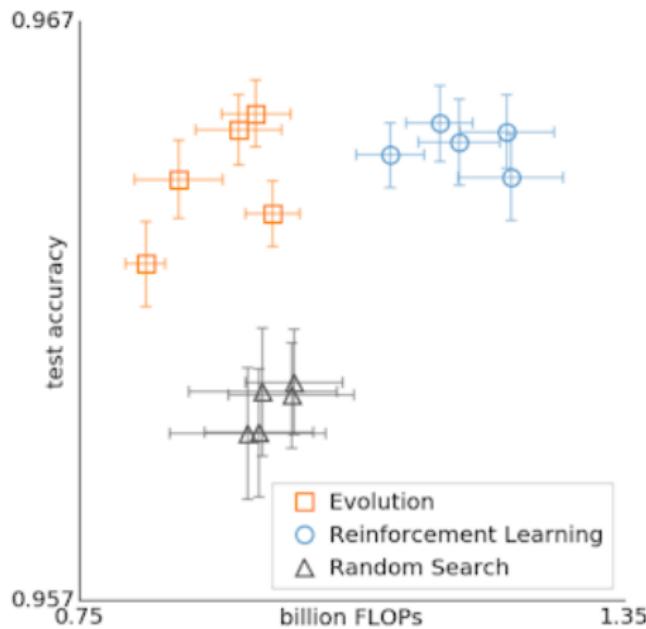
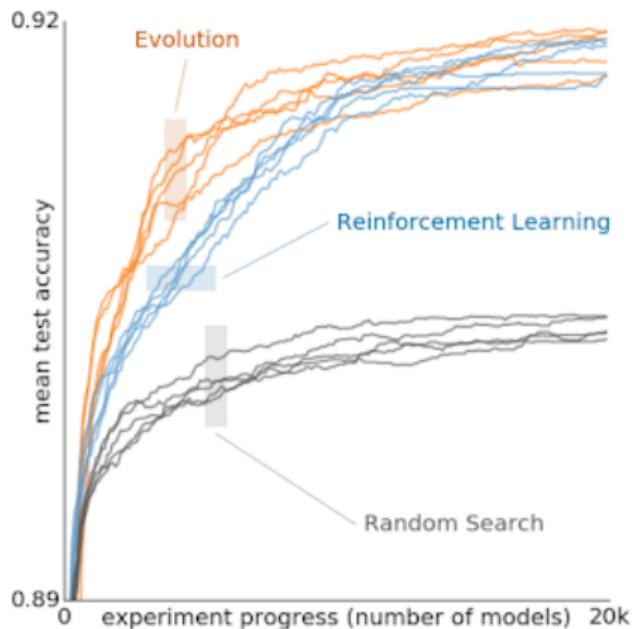


Convolutional Neural Fabrics

- problem: how to find the most relevant architecture
- todays solution: try and test
- A new solution: learn the architecture

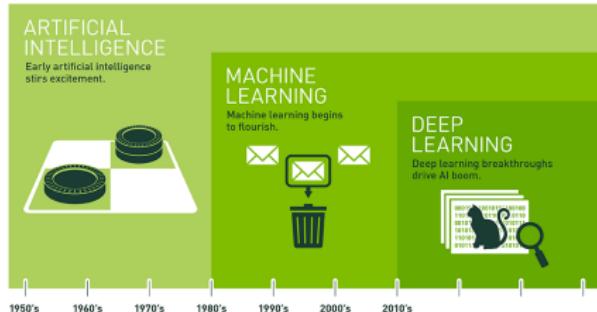


Neural Architecture Search



Road map

- 1 Why deep learning?
- 2 The first stage: 1890 - 1969
- 3 The second stage: 1985 - 1995
- 4 The third stage: 2006 - (2012) - 2018...
- 5 What's new in deep learning?
 - Big is beautiful
 - Two Hot topics: data and architecture
- 6 Conclusion



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

The deep learning time line

- The first stage: 1890 - 1969
 - ▶ learning is optimization with **stochastic gradient** (to scale)
- The second stage: 1985 - 1995
 - ▶ NN are universal approximator **differentiable graphs** (that scales)
- The third stage: 2006 - (2012) - 2018...
 - ▶ scale with **big** data+computers+architecture (deep)
- Open issues
 - ▶ Many...

Open issues in Deep learning : A critical appraisal

arXiv.org > cs > arXiv:1801.00631

Search or Art

(Help | Advanced)

Computer Science > Artificial Intelligence

Deep Learning: A Critical Appraisal

Gary Marcus

(Submitted on 2 Jan 2018)

Although deep learning has historical roots going back decades, neither the term "deep learning" nor the approach was popular just over five years ago, when the field was reigned by papers such as Krizhevsky, Sutskever and Hinton's now classic (2012) deep network model of Imagenet. What has the field discovered in the five subsequent years? Against a background of considerable progress in areas such as speech recognition, image recognition, and game playing, and considerable enthusiasm in the popular press, I present ten concerns for deep learning, and suggest that deep learning must be supplemented by other techniques if we are to reach artificial general intelligence.

Comments: 1 figure

Subjects: Artificial Intelligence (cs.AI); Machine Learning (cs.LG); Machine Learning (stat.ML)

MSC classes: 97R40

ACM classes: I.2.0; I.2.6

Cite as: arXiv:1801.00631 [cs.AI]

(or arXiv:1801.00631v1 [cs.AI] for this version)

Bibliographic data

[Enable Bibex (What is Bibex?)]

Submission history

From: Gary Marcus [view email]

[v1] Tue, 2 Jan 2018 12:49:35 GMT (258kb)

For most problems where deep learning has enabled transformationally better solutions (vision, speech), we've entered diminishing returns territory in 2016-2017.

Francois Chollet, Google, author of Keras neural network library Dec. 2017

10 Limits on the scope of deep learning

- It is data hungry
- specialized training (no learning)
- it has no natural way to deal with hierarchical structure
- it has struggled with open-ended inference
- it is not sufficiently transparent
- not well integrated with prior knowledge (NLP)
- cannot distinguish causation from correlation
- assume stationarity
- can be fooled (it is not robust)
- is difficult to engineer with

10 Limits on the scope of deep learning

GREEN LEARNNING

- It is data hungry
- specialized training (no learning)
- it has no natural way to deal with hierarchical structure
- it has struggled with open-ended inference

THEORY NEEDED

- it is not sufficiently transparent

- not well integrated with prior knowledge (NLP)

- cannot distinguish causation from correlation

- assume stationarity

ADVERSARIAL TRAINING

- can be fooled (it is not robust)

- is difficult to engineer with

AUTO ML

To go further

- books
 - ▶ I. Goodfellow, Y. Bengio & A. Courville, *Deep Learning*, MIT Press book, 2016
<http://www.deeplearningbook.org/>
 - ▶ Gitbook leonardoaraujosantos.gitbooks.io/artificial-intelligence/
- conferences
 - ▶ NIPS, ICLR, xCML, AIStats,
- Journals
 - ▶ JMLR, Machine Learning, Foundations and Trends in Machine Learning, machine learning survey <http://www.ml-surveys.com/>
- lectures
 - ▶ Deep Learning: Course by Yann LeCun at Collège de France in 2016
college-de-france.fr/site/en-yann-lecun/inaugural-lecture-2016-02-04-18h00.htm
 - ▶ Convolutional Neural Networks for Visual Recognition (Stanford)
 - ▶ deep mind (<https://deepmind.com/blog/>)
 - ▶ CS 229: Machine Learning at stanford Andrew Ng
- Blogs
 - ▶ Andrej Karpathy blog (<http://karpathy.github.io/>)
 - ▶ <http://deeplearning.net/blog/>
 - ▶ <https://computervisionblog.wordpress.com/category/computer-vision/>